

Phase 4

Runtime & Self-Hosting

Bytecode VM, JIT Compilation via Cranelift, and the first self-hosted component — a lexer written in Metalogos.

Date: 2026-06-01

Sub-phases: 4.1 / 4.2 / 4.3 / 4.4

Tests: 25/25 GREEN

ADRs: 0020-0023




Table of Contents

1. Introduction	2
1.1 Phase 4 Sub-phases	2
2. Phase 4.1: Bytecode Instruction Set + Stack VM	3
2.1 Architecture	3
2.2 Instruction Set	3
2.3 METALOGOS-Specific Opcodes	4
2.4 Files Created	4
3. Phase 4.2: VM Full Feature Coverage	5
3.1 Rule Priority Sorting	5
3.2 Collections Flag Propagation	5
3.3 Performance Benchmarks	5
3.4 Strict Dual-Mode Testing	6
4. Phase 4.3: JIT Compilation via Cranelift	6
4.1 Architecture	7
4.2 Purity Analysis	7
4.3 Intentional Constraints	7
4.4 Performance Analysis	8
5. Phase 4.4: Self-Hosting (Polar Star)	9
5.1 Hybrid Architecture	9
5.2 Metalogos Seven Pillars Usage	9
5.3 New Builtins	10
5.4 Architecture Changes	10
5.5 Contract Verification	10
6. Phase 4 Summary	11

1. Introduction

This report summarizes the complete implementation of **Phase 4** of the METALOGOS programming language project. Phase 4 is the final milestone of the M1 roadmap, transforming METALOGOS from a pure tree-walking interpreter into a multi-backend runtime system with bytecode compilation, stack-based virtual machine execution, JIT compilation for hot paths, and the first self-hosted component written in Metalogos itself.

METALOGOS is an AI-native programming language built on seven pillars: Entity, Pattern, Flow, Memory, Rule, Learn, and Adapt. Phases 1-3 established the core language features: a pest PEG parser producing an AST, semantic analysis with diagnostics, and a tree-walking interpreter supporting all seven pillars with Fluid Types, confidence propagation, entity store, type inference, knowledge graph relations, LLM-backed learnable patterns, and adaptive few-shot mutation.

Phase 4 consists of four sub-phases, each delivered as a formal work order (Naryad). The sub-phases are incrementally built on top of each other, starting from the foundational bytecode VM and culminating in a self-hosting proof-of-concept that demonstrates Metalogos can process its own source code. All 25 tests pass across both the working workspace and the GitHub repository.

1.1 Phase 4 Sub-phases

Sub-phase	Name	Description	Tests
4.1	Bytecode VM	Stack-based bytecode + compiler + interpreter	15
4.2	Full Coverage	Rule priority, collections, benchmarks	18

4.3	JIT Hot Paths	Cranelift JIT for pure Float patterns	24
4.4	Self-hosting	Lexer written in Metalogos	25

Table 1: Phase 4 sub-phases and cumulative test counts

2. Phase 4.1: Bytecode Instruction Set + Stack VM

Phase 4.1 introduced a bytecode compilation layer and stack-based virtual machine, providing a compilation target separate from the AST. This enables future optimization passes and eliminates the overhead of repeated AST traversal inherent in tree-walking interpretation. The implementation consists of three new modules forming the compilation and execution pipeline.

2.1 Architecture

The bytecode pipeline follows a clean three-stage architecture: AST declarations are compiled into a Program data structure by `compiler.rs`, which is then executed by `vm.rs`. The Program contains global variable slots, compiled pattern functions, learnable pattern metadata, rule definitions, and the main instruction sequence. The compiler operates in two passes: the first pass collects type information and assigns global variable slots, while the second pass generates bytecode instructions with context-aware expression compilation that distinguishes between pattern parameters (compiled as `LoadLocal`) and global variables (compiled as `LoadGlobal/LoadGlobalByName`).

2.2 Instruction Set

The VM implements 30+ opcodes organized into eight functional categories. Each opcode is a variant of the Instruction enum in `bytecode.rs`, with operands embedded

inline. The instruction set includes stack manipulation (Const, LoadGlobal,StoreGlobal, LoadLocal), function operations (RegisterPattern, CallBuiltin,CallPattern, Return), arithmetic (Add, Sub, Mul, Div), comparison (Contains,CmpGt, CmpLt, CmpGe, CmpLe, CmpEq), struct operations (MakeStruct, GetField),Fluid type operations (MakeFluid, Collapse), control flow (Jump, JumpIfNot,JumpIfLow), METALOGOS-specific memory operations (Memorize, Recall, Forget), LLMdispatch (LlmCall), and pipeline macro instructions (FlowExec, ExecuteRules).

2.3 METALOGOS-Specific Opcodes

Three opcodes are unique to METALOGOS and have no equivalent in conventional VMs. **JumpIfLow** implements confidence-based branching: it pops a value and jumps if the value represents a confidence score below a given threshold, enabling flow branch syntax. **Collapse** collapses a Fluid superposition to a concrete type by finding the highest-confidence variant matching the required type name, returning Unit if below the collapse threshold of 0.1. **Memorize/Recall/Forget** implement the memory pillar with priority-based storage, substring-based recall with exponential decay, and time-based forgetting. These operations are first-class VM instructions rather than library calls, ensuring consistent semantics across all execution backends.

2.4 Files Created

File	Line s	Purpose
src/bytecode.rs	247	Instruction enum, Program, CompiledFn, CallFrame
src/compiler.rs	~470	AST to bytecode: two-pass compilation, purity analysis
src/vm.rs	~950	Stack VM: dispatch loop, pattern calls, rule engine
src/lib.rs	modi f.	Added VM modules, run_program_vm functions
docs/adr/0020-bytecode-vm.md	103	Architecture Decision Record

Table 2: Files created/modified in Phase 4.1

All 11 golden test examples produce identical output through both the tree-walking interpreter and the bytecode VM. Zero warnings, zero clippy errors. The implementation achieved 100% feature parity across all METALOGOS constructs including struct types, fluid types, rules with priority, learnable patterns with LLM dispatch, memory operations with decay, knowledge graph relations, adapt/mutate operations, import resolution, and flow pipelines with branching.

3. Phase 4.2: VM Full Feature Coverage

Phase 4.2 addressed four semantic gaps identified between the VM and tree-walking interpreter: rule priority sorting, collections flag propagation, lenient testing, and the absence of a performance baseline. Each gap was systematically resolved with targeted code changes.

3.1 Rule Priority Sorting

The interpreter sorts rules by priority descending before execution using a stable sort that preserves declaration order for ties. The VM was evaluating rules in declaration order, which was correct for single-rule programs but semantically incorrect for multi-rule programs with different priorities. The fix added a priority field to `CompiledRule` in `bytecode.rs` and implemented the same stable descending sort in `vm.rs`, ensuring identical execution semantics.

3.2 Collections Flag Propagation

The compiler tracked the `collections_loaded` flag but never passed it to the VM `Program` struct, meaning `map/filter/reduce` would never be recognized as valid operations in the VM path. The fix added a `collections_loaded` field to `Program` and ensured the compiler propagates this flag through the compilation pipeline.

3.3 Performance Benchmarks

A comprehensive benchmark suite was added in `tests/bench.rs` with synthetic programs ranging from 10 to 1000 chained pattern calls. The VM provides a modest but consistent speedup that scales with program complexity.

Steps	Tree-Walking	VM	Speedup
10	170 us	155 us	1.10x
50	340 us	324 us	1.05x
100	565 us	499 us	1.13x
500	2516 us	1988 us	1.27x
1000	4520 us	3529 us	1.28x

Table 3: VM vs Tree-Walking benchmark results (dev profile)

The speedup increases with pipeline length (1.05x at 50 steps to 1.28x at 1000 steps), indicating that the VM advantage grows with program complexity. The aggregate speedup across all golden examples is 1.17x (TW: 5601 us vs VM: 4769 us). The modest speedup is expected because both paths use the same `Value` type with `clone()` on every stack operation, `FlowExec` remains a macro instruction, and pattern body execution uses a separate dispatch loop rather than inlining.

3.4 Strict Dual-Mode Testing

A new test `all_vm_examples_match_tree_walking` was added to `tests/vm_golden.rs`. This test runs every `.mlog` example through both tree-walking and VM, then performs `assert_eq` on the trimmed output. Any divergence causes an immediate test failure with a diff showing TW vs VM output. This provides a hard contract guarantee that VM output equals tree-walking output for all examples.

4. Phase 4.3: JIT Compilation via Cranelift

Phase 4.3 introduced a JIT compilation layer that compiles frequently-called pure Float-only patterns to native machine code via Cranelift, bypassing the bytecode

dispatch loop entirely. The JIT is an optional layer on top of the VM, enabled via the `--jit` CLI flag.

4.1 Architecture

The JIT pipeline follows a hot-path detection strategy: the VM counts calls per pattern, and when a pure Float-only pattern exceeds the threshold (default $N=100$ calls), it is compiled to native code via Cranelift. Subsequent calls dispatch directly to the compiled function pointer, bypassing the bytecode dispatch loop. The new `jit.rs` module wraps Cranelift `JITModule` and provides four key operations: `is_jit_eligible` checks if a pattern can be JIT-compiled, `compile_pattern` translates bytecode to Cranelift IR and emits native code, and `call_jit` dispatches to the native code via arity-specific transmuters supporting 0-3 arguments.

4.2 Purity Analysis

The compiler marks patterns as `is_pure` via `analyze_purity()`. A pattern is pure if all parameters are Float, and the body contains only `LoadLocal`, `Const(Float)`, `Add`, `Sub`, `Mul`, `Div`, `CmpGt/CmpLt/CmpGe/CmpLe/CmpEq`, and `Return`. JIT eligibility is a stricter subset that additionally excludes comparison instructions because Cranelift boolean-to-float conversion requires instructions not available in the current version. Patterns with globals, builtins, LLM calls, struct operations, or memory access are never JIT-eligible.

4.3 Intentional Constraints

Constraint	Rationale
Float-only parameters	Cranelift cannot safely interact with Rust String type (no GC integration)
Arithmetic-only bodies	Comparison boolean-to-float conversion adds complexity without proportional benefit
No globals access	Eliminates need for guard checks (no invalidation tracking)
No side effects	Ensures deterministic behavior for native code
Learnable/adapt stays in VM	LLM interactions and adapt mutations are non-deterministic with I/O side effects
No guard checks	Pure patterns have no external dependencies, always produce correct results

Table 4: JIT constraints and rationale

4.4 Performance Analysis

Steps	VM-only	VM+JIT	JIT Speedup
10	22 us	91 us	0.24x
50	42 us	105 us	0.40x
100	62 us	172 us	0.36x
500	246 us	291 us	0.85x
1000	482 us	505 us	0.95x

Table 5: JIT benchmark results (release mode, optimized)

The benchmark shows JIT is slower than VM-only for small programs but approaches parity at scale. The fixed JIT overhead of approximately 70 microseconds includes JITModule creation (ISA detection, memory allocation for executable pages), Cranelift compilation of the first hot pattern, and `finalize_definitions()` to make code executable. At 1000 steps, the per-step cost converges: VM is 0.48 us/step while JIT achieves 0.44 us/step, making JIT approximately 8% faster per-step at scale. The crossover point where JIT breaks even is approximately 1500-2000 pattern calls.

All 13 golden test examples produce identical output across three execution modes: tree-walking interpreter, bytecode VM (`mlog run --vm`), and bytecode VM + JIT (`mlog run --jit`). The CLI supports three execution modes with flexible threshold configuration.

Command	Backend	Description
<code>mlog run file.mlog</code>	Tree-walking	Default AST interpreter
<code>mlog run --vm file.mlog</code>	Bytecode VM	Stack-based bytecode execution
<code>mlog run --jit file.mlog</code>	VM + JIT	Hot pure patterns compiled to native code
<code>mlog run --jit --jit-threshold 50 file.mlog</code>	VM + JIT	Custom hot-path threshold

Table 6: CLI execution modes

5. Phase 4.4: Self-Hosting (Polar Star)

Phase 4.4 is the final sub-phase and delivers the first self-hosted component: a lexer written in Metalogos that processes Metalogos source code. This demonstrates that the language has sufficient expressiveness to reason about its own syntax, a fundamental milestone for any programming language project. The implementation uses a hybrid architecture where Rust handles low-level character scanning and Metalogos handles high-level token classification.

5.1 Hybrid Architecture

The self-hosted lexer divides responsibilities between Rust and Metalogos intentionally. Metalogos has no character-level iteration or loop constructs, making character scanning inherently sequential and better suited to Rust. The linguistic analysis, including keyword recognition and type classification, is done in Metalogos, demonstrating the language symbolic processing capability through its core pillars: patterns for token type recognition, flow for the tokenization pipeline, and memory for the keyword lookup table.

5.2 Metalogos Seven Pillars Usage

Pillar	Usage in Lexer
Entity	Global variable "source" holds stdin input
Pattern	IsKeyword, IsFloatLit, IsOperator, IsStringLit, ClassifyToken, Tokenize, Classify, Format
Flow	Lexer pipeline: source - Tokenize - Classify - Format - output
Memory	34 keywords stored via memorize, looked up via recall
Rule	Not directly used (lexer is purely pattern-based)
Learnable	Not used (lexer is deterministic)
Adapt	Not used (no training needed for token classification)

Table 7: Metalogos seven pillars usage in the self-hosted lexer

5.3 New Builtins

Six new builtins were added to the runtime to enable self-hosting. The `stdin()` builtin reads all `stdin` into a `String`, providing the input source for the lexer. `split_tokens(s)` performs character-level tokenization respecting string literals and multi-char operators (`-`, `==`, `=`, `=`), delegating sequential character iteration to Rust. The `if_eq(a, b, then, else)` builtin provides ternary conditional for expressions supporting both `String` and `Float` equality. `newline()` returns a newline character for output formatting. `is_string_token(s)` detects string literals by checking start/end quotes.

5.4 Architecture Changes

Two significant architecture changes were made to support self-hosting. First, `map/filter/reduce` collection operations are now always available without imports. Previously gated behind `"import std/collections"` which required `stdlib` to be in a specific path relative to the source file, this change enables self-hosting programs in any directory. Second, the collection functions were refactored to use a new `eval_collection_fn` helper that directly evaluates pattern bodies without going through `invoke()`, avoiding a mutable borrow issue when called from `eval_expr_with_env`.

5.5 Contract Verification

The contract specifies that running the self-hosted lexer on the hello world example should produce a specific token list. The verification command is:

```
$ mlog run self-host/lexer.mlog < examples/m1_hello.mlog
```

The output correctly classifies all 37 tokens into `KEYWORD` (`entity`, `pattern`, `flow`, `return`, `input`, `output`, `s`), `IDENT` (`greeting`, `Shout`, `String`, `Main`, `upper`), `OPERATOR` (`:`, `=`, `(`, `)`, `-`, `+`, `{`, `}`), and `STRING` (`"Hello, Metalogos!"`, `"!"`) categories. The integration test in `tests/self_host_lexer.rs` spawns the `mlog` process with piped `stdin` and asserts exact match with the expected output file.

6. Phase 4 Summary

Phase 4 transformed METALOGOS from a tree-walking interpreter into a multi-backend runtime system. The four sub-phases were delivered incrementally, each building on the previous, with all changes validated by automated tests that guarantee semantic equivalence across all execution backends.

Metric	Value
Total sub-phases	4 (4.1 - 4.4)
ADRs published	4 (0020 - 0023)
New Rust modules	3 (bytecode.rs, compiler.rs, jit.rs)
New .mlog programs	1 (self-host/lexer.mlog)
New builtins	6 (stdin, split_tokens, if_eq, newline, is_string_token, len ext.)
VM opcodes	30+
Execution backends	3 (Tree-walking, VM, VM+JIT)
Total tests	25 (all green)
Golden examples	13 (all pass in all backends)
VM speedup over TW	1.17x aggregate, up to 1.28x at scale
JIT speedup over VM	~8% per-step at 1000+ iterations
Self-hosting	Lexer processes Metalogos source code

Table 8: Phase 4 summary metrics

Phase 4 is closed. METALOGOS has a bytecode VM with full language feature coverage, JIT compilation for hot pure patterns via Cranelift, and the first self-hosted component. The M1 roadmap through Phase 4 is complete. The project is ready for M2 milestones which will extend the language with further self-hosting

components, optimization passes, and advanced runtime features.